


Java to C++ (for NumCS)

1. Core Language

Felix Friedrich
Malte Schwerhoff
D-INFK 2023

This Course



- Goal: ease transition from Java to C++
- Expects initial knowledge in Java
- Tailored to NumCS – not a generic C++ course
- Not self-contained: the gist rather than the details
- Second iteration: **Please interact, ask, and give feedback**

Who Is Involved?



Dr. Felix Friedrich



Dr. Malte Schwerhoff

- D-INFK service lecturers
- Experienced C++ lecturers
- Lecturers first two weeks






Dr. Vasile Gradinaru

←
*picture maybe
somewhat dated*

- D-MATH (CSE) lecturer
- Experienced numerical methods lecturer
- Main lecturer
(after 2nd week)

Good Sources of Information (on C++ and Eigen)

- <https://cppreference.com>
- Programming: Principles and Practice Using C++
(Stroustrup; <https://u.ethz.ch/2MgzB+>)
- A Tour of C++
(Stroustrup; <https://u.ethz.ch/aeH7d+>)
- Effective Modern C++
(Scott Meyers; <https://u.ethz.ch/m1DvT+>)
- Eigen (linear algebra library): <https://eigen.tuxfamily.org/dox/>



A Bit of Background on C++

Why C++






- Important language in **computational science, engineering, high-performance computing**
- NumCS uses C++ to discuss implementations of numerical methods
- Often combined with Python (not this course): user-friendly Python interface to a high-performance C++ implementation



**401-0663-00L Numerical Methods
for Computer Science**

What C++ Is Like

- C++ has a certain reputation ... partly deserved
- But **modern C++** (\geq C++11, latest standard is C++20) has come a long way



What C++ Is Like



- C++ is not exactly beginner-friendly
 - Hardware and OS “shine through” — great for experts, hard for beginners
 - Many choices — great for experts, hard for beginners
 - Explicit choices — explicate decisions, at cost of syntactical convenience



- But C++ is also not opinionated. Two examples:
 - Java-style OOP with **runtime polymorphism**; but also **compile-time polymorphism** and static meta-programming
 - **Value semantics** (think Java primitive types); but also **reference semantics** (think Java objects)

C++ vs. Java – Big-Picture Differences



Compiled to **machine code**

- Binaries not portable (OS, hardware)
- Static code optimisations



Emphasises **efficiency** (runtime, memory)



Extended C

- Inherited efficiency emphasis
... and some design problems
- Largely interoperable with C



Compiled to **intermediate code**

- Portable, but requires
intermediate layer (JVM)
- Runtime code optimisations



Emphasises **safety**



Designed from scratch, with C/C++ lessons-learned in mind

What Does C++ Code Look Like?

```
#include <iostream>
#include <limits>

double average(std::istream&);

int main() {
    std::cout << "Enter values to average: ";
    double avg = average(std::cin);
    std::cout << "Average = " << avg << "\n";
}

// Returns the average of numbers read from the given input stream
double average(std::istream& in) {
    double value, sum = 0; // Attention: value is not initialised
    unsigned count = 0;

    while (in >> value) {
        ++count;
        sum += value;
    }

    return count != 0 ? sum / count : std::numeric_limits<double>::signaling_NaN();
}
```

Java vs. C++: Some Similarities and Differences

Some Similarities between C++ and Java

- Core syntax
 - Case sensitivity, curly braces, parentheses, semicolon, comments
- Operators
 - `+`, `*`, `=`, `+=`, `==`, `++`, `<`, `&&`, ...
 - Same syntax, precedence, short-circuiting, mixed types conversion
- Fundamental types: `int`, `float`, `double`, `char`, `bool`; also `void`
- Control structures: `return`, `if-else`, `switch`, `for`, `while`,
- Code modularity
 - Functions, classes
 - Namespaces (and *modules* with C++20)

Differences: Fundamental Types in C++

- Signed and unsigned integral types
 - E.g. `int` vs. `unsigned int` (or just `unsigned`)
- Width of `bool`, `char`, `int`, `long` is **implementation-defined**, with lower bounds
 - E.g. an `int` is at least 16 bits wide, a `long` 32 bits
 - Fixed-width types introduced with C++11, e.g. `int32_t`
- <https://en.cppreference.com/w/cpp/language/types>
- <https://en.cppreference.com/w/cpp/types/integer>

OS	Architecture	Size of "long" type
Windows	IA-32	4 bytes
	Intel® 64	4 bytes
Linux	IA-32	4 bytes
	Intel® 64	8 bytes
mac OS	Intel® 64	8 bytes

Size of 'long integer' data type (C++) on various architectures and OS ([intel.com](https://en.cppreference.com/w/cpp/types/integer))

The Price of Efficiency and Generality

- Background

1. C++ aims to enable **efficient programs** on **any hardware & OS**
2. C++ is old (limited language design experience, fewer hardware & OS standards, significantly fewer computational resources)

- Intuitive trade-off: rigidly defined language semantics increase safety, but complicate efficient hardware & OS support

- C++ language standard grants freedom to language implementers

1. **Implementation-defined behaviour:** behaviour may vary, must be documented
2. **Unspecified behaviour:** implementation-defined, no need to document
3. **Undefined behaviour:** if present, program is not required to do anything meaningful

Common Sources of Undefined Behaviour

- *Signed* integer over-/underflow
 - Java: well-defined
- Out-of-bounds array access, e.g. `data[data.size()]`
 - Java: runtime check, exception
 - C++: `data.at(idx)` performs runtime check
- Null-pointer dereferencing
 - Java: runtime check, exception
- Read & write *same memory location* in single expression, e.g. `cout << (x + x++)`
 - Java: well-defined
- Reading uninitialized variables
 - Java: compiler check for local variables, default initialisation for member variables
- Accessing destructed objects/deallocated memory
 - Java: cannot happen, since garbage collector manages memory



Differences: Initialisation

- C++ default-initialises variables iff the type declares a default constructor
 - Applies to local variables, and member variables (object fields)

- No default constructors for fundamental types (int, double, ...)

```
int count; // Not initialised  
cout << count; // Undefined behaviour
```

C++

```
int count; // Not initialised  
println(count); // Compiler error
```

Java

- C++: reading uninitialised variables is undefined behaviour
- Java:
 - Guaranteed compiler error for local variables
 - Default initialisation for member variables


Undefined Behaviour

- Compile- or runtime checks might be performed, but **no guarantees**
- **Keep existence in mind**
 - To avoid common problems (previous slides)
 - During debugging (undefined behaviour often results in nondeterministic behaviour)
- **Otherwise not central for this course**
 - We will not go into further details
 - Newer C++ versions sometimes replace *undefined* by *unspecified behaviour*
- If interested: <https://en.cppreference.com/w/cpp/language/ub>
- Demo of what *may* happen: <https://gcc.godbolt.org/z/MsjPnWs5Y>

Differences: Expression Evaluation

HANDOUT SLIDE

- Evaluation order of subexpressions
 - *Unspecified behaviour* in C++
 - Specified as left to right in Java
- Compound expression example: $e1 + e2 * e3$
 - Java: $e1$ evaluated first, then $e2$, finally $e3$
 - C++: Any order allowed
- Likewise for function call arguments, e.g. $\text{fun}(e1, e2, e3)$
- Relevant (only) when expressions have side-effects



Differences: Automatic Numerical Conversions

HANDOUT SLIDE

- Automatic conversion to `bool`
 - In C++, variables of fundamental type can be used as `bool` without explicit casting
- Automatic narrowing between numeric types
 - In C++, e.g. assigning a `double` to an `int` does not require an explicit cast
- Automatic numerical promotions in mixed numerical expressions
 - E.g. `some_short * some_double + some_int` → all values promoted to doubles
 - Rules analogous to Java's:
https://en.cppreference.com/w/cpp/language/operator_arithmetic#Conversions
 - C++-specific: signed promoted to unsigned → `-1 < 1u` is false
<https://www.modernescpp.com/index.php/safe-comparisons-of-integrals-with-c-20>

Differences: Type Inference

- C++ is statically typed, just like Java
- Limited support for **type inference** includes
 - Type parameter interference → templates, 2nd week
 - Keyword **auto**

```
std::vector<int> data = ...;  
std::vector<int>::iterator head = data.begin();
```

vs.

```
std::vector<int> data = ...;  
auto head = data.begin();
```

- Types serve as code documentation → don't overuse **auto**

```
auto model = extract_model(...); // What's model? Which functions does it provide?
```

Values vs. References

Value Types / Pass by Value

- Consider the example on the right
- Can you explain the observed behaviour?

```
void swap(int x, int y) {  
    int tmp = x;  
    x = y;  
    y = tmp;  
}
```

```
int main() {  
    int x = 1;  
    int y = 2;  
    swap(x, y);  
    std::cout << x << ", " << y; // 1, 2  
}
```

just as
in Java

Value Types / Pass by Value

- Consider the extended example. Can you explain the observed behaviour?

```
struct Coordinate { // Java class
    int x;
    int y;

    Coordinate(int a, int b) {
        x = a;
        y = b;
    }
};
```

```
void swap(Coordinate c) {
    int tmp = c.x;
    c.x = c.y;
    c.y = tmp;
}

int main() {
    Coordinate c(1, 2); // Obj. construction
    swap(c);
    std::cout << c.x << ", " << c.y; // 1, 2
}
```

*not as
in Java*

Value Types / Pass by Value

- By default, types in C++ are **value types**
 - Fundamental types (**int**, etc.): *same* as in Java
 - Instances of classes: *different* from Java
- Value types (of statically-known size) can be allocated on the stack
 - Typically faster than heap accesses
- Pass by value: values are **copied** upon function calls
- Return by value: analogous

Outlook: Values vs. Pointers

- Java's *references* are (a simplified version of) *pointers* in C++

```
Coordinate* location = new Coordinate(1, 2);
```

- The pointer value (memory address) is still passed by value (i.e. copied), as in Java
- C++ has no garbage collector → manual memory management
 - Every **new** needs a **delete**
 - Danger of memory leaks, accessing deleted objects, freeing memory twice
 - *Smart pointers* simplify memory management (while preserving efficiency)
- Pointers are not relevant for NumCS → not discussed further

Pass by Reference / Reference Types


- Pass by value not always desirable
 1. Effects: modifying a shared object (e.g. a shared counter in a concurrent system)
 2. Efficiency: avoiding copying large data (e.g. a huge matrix)
- Values can be shared via C++ references
 - Still no heap involved (efficiency)

Pass by Reference / Reference Types

- Pass-by-reference semantics with **C++ references**

```
struct Coordinate {  
    ... // same as before  
};
```

```
void swap(Coordinate& c) {  
    int tmp = c.x;  
    c.x = c.y;  
    c.y = tmp;  
}  
  
int main() {  
    Coordinate c(1, 2);  
    swap(c);  
    std::cout << c.x << ", " << c.y; // 2, 1  
}
```



just as
in Java

Pass by Reference / Reference Types

- References *can* (nearly) be used everywhere, e.g.

```
int anakin_skywalker = 7; // Happy birthday!  
int& darth_vader = anakin_skywalker;  
darth_vader = 40; // Oh oh, Anakin rapidly aged ...
```

- Reference types are *typically* used
 - As function parameters
 - In lambda expressions (2nd week)
 - As member variables (classes, this Thursday)

Pass by Reference / Reference Types

- References must be initialised

```
int& darth_vader; // Compiler error
```

- References cannot be re-aliased:

```
int& darth_vader = anakin_skywalker;  
darth_vader = kermit_the_frog;  
// vader (and thus anakin) gets kermit's value,  
// but vader does not become an alias of kermit
```

Pass by Reference / Reference Types – With Fundamental Types

```
void swap(int& x, int& y) {  
    int tmp = x;  
    x = y;  
    y = tmp;  
}  
  
int main() {  
    int x = 1;  
    int y = 2;  
    swap(x, y);  
    std::cout << x << ", " << y; // 2, 1  
}
```

- Can you do this in Java?

Constants

- Java's `final` is `const` in C++

```
const unsigned speed_of_light = 299792458;  
...  
speed_of_light *= 2; // Compiler error
```

- A program is `const-correct` if all variables not intended to be mutated are typed as `const`
- Concept lifted to *constant expressions* (= evaluable at runtime) to support compile-time meta-programming
 - No further details here
 - If interested: https://en.cppreference.com/w/cpp/language/constant_expression

Const-References

- Consider the code below. Can it be improved?

```
bool is_sorted(std::vector<int> v) {  
    for (unsigned i = 1; i < v.size(); ++i) {  
        if (v[i] < v[i-1])  
            return false;  
    }  
  
    return true;  
}  
  
std::vector<int> v(10'000'000); // Size 10e7  
...  
if (is_sorted(v)) ...
```


Const-References

- Yes: **const-reference** for efficiency (no copy) and safety (no mutation)

```
bool is_sorted(const std::vector<int>& v) {  
    // ... as before ...  
}  
  
std::vector<int> v(10'000'000);  
...  
if (is_sorted(v)) ...
```

Functions & Operators

Functions

- Core syntax as in Java: $R \text{ name}(T1 \ p1, \dots, Tn \ pn) \{ \text{body} \}$

- Function overloading

```
void process(int data) { ... }  
void process(double data) { ... }  
  
int x = ...;  
  
process(x); // calls 1st function  
process(3.13); // calls 2nd function
```

- Default parameters

```
// Returns true if data contains elem  
// at least rep times  
bool occurs(  
    int elem,  
    std::vector<int> data,  
    unsigned rep = 1) {  
  
    ...  
}  
  
std::cout << occurs(0, my_data);  
std::cout << occurs(77, my_data, 2);
```

Functions

HANDOUT SLIDE

- Variadic functions

```
double average(double values...) {  
    // compute average of given values  
}  
  
double avg1 = average(2.5, 1.7);  
double avg2 = average(avg1, 13, 7.4, 28.9);
```

- Declaration very similar to Java – looks nice ...
- ... but usage inside the function does not, thus omitted
- <https://en.cppreference.com/w/cpp/utility/variadic>
- Alternative: variadic *templates* (2nd week)

Operator Overloading

- In contrast to Java, C++ supports **overloading of operators**

```
struct Rational {  
    int n; int d;  
  
    Rational(int _n, int _d) { n = _n; d = _d; }  
};
```

```
Rational operator+(const Rational& l, const Rational& r) {  
    return Rational(l.n * r.d + l.d * r.n, l.d * r.d);  
}
```

```
std::ostream& operator<<(std::ostream out&, const Rational& r) {  
    out << r.n << ", " << r.d;  
    return out; // enables chaining of <<  
}
```

```
Rational r1(1, 3);  
Rational r2(5, 7);
```

```
std::cout << (r1 + r2);
```

- <https://en.cppreference.com/w/cpp/language/operators>

Separate Compilation

Semantics of #include


main.cpp

```
#include <rational.cpp>
```

```
int main() {  
    Rational r1 = ...;  
    ...  
}
```



compile
main.cpp



main.cpp (as compiled)

```
struct Rational {  
    int n;  
    int d;  
  
    Rational(int a, int b) {  
        n = a;  
        d = b;  
    }  
};  
  
...  
  
int main() {  
    Rational r1 = ...;  
    ...  
}
```

rational.cpp

```
struct Rational {  
    int n;  
    int d;  
  
    Rational(int a, int b) {  
        n = a;  
        d = b;  
    }  
};  
  
...
```

- `#include <file.cpp>` is replaced with the content of `file.cpp`
- Possible disadvantages?

Includes and Separate Compilation

- Fact: `#include <lib.cpp>` is replaced with the contents of file `lib.cpp`
- New goal: Enable separate compilation of units of code (e.g. classes, libraries); resulting machine code then linked afterwards
 - Improves performance (shared library only compiled once)
 - Allows binary distributions (e.g. closed-source libraries)
- Question: Given how `#include` works, how to enable separate compilation?
- C++ solution: Separation of declarations from definitions in separate files
 - Header files `.h` contain declarations, implementations reside in `.cpp` files

Header vs. Implementation

main.cpp

```
#include <rat.h>

int main() {
    Rat r1 = ...;
    ...
}
```

rat.h

```
struct Rat {
    int n;
    int d;

    Rat(int a, int b);
};

Rat operator+(
    const Rat& l,
    const Rat& r);

// more declarations ...
```

rat.cpp

```
Rat::Rat(int a, int b) {
    n = a;
    d = b;
}

Rat operator+(const Rat& l, const Rat& r) {
    return Rat(l.n * r.d + l.d * r.n, l.d * r.d);
}

// more implementations ...
```

- Declarations suffice for compilation
- **rat.cpp** and **main.cpp** can be compiled separately

```
$ g++ -c main.cpp    (produces main.o)
$ g++ -c rat.cpp      (produces rat.o)
```
- ... and linked together later

```
$ g++ rat.o main.o -o main    (produces executable main)
```

Preventing Re-inclusion

- If `rat.h` is (transitively) included multiple times, symbols are **redeclared**
- Avoid by using *pre-processor macros* (not discussed further) to establish **include guards**:

rat.h

```
#ifndef RAT_H
#define RAT_H

struct Rat {
    int n;
    int d;

    Rat(int a, int b);
};

// more declarations ...

#endif // RAT.H
```

- Alternative using *preprocessor directive* (not discussed further)
 - Certain advantages over include guards
 - Not supported by all compilers

rat.h

```
#pragma once

struct Rat {
    int n;
    int d;

    Rat(int a, int b);
};

// more declarations ...
```

Separate Compilation on Code Expert

- Code Expert environments typically don't give you access to the compiler command-line
 - Advantage: simpler to use, less potential for mistakes
 - Disadvantage: separate compilation might not be possible
- Probably only relevant if you create your own files
(in contrast to completing given skeleton files)